

Average execution times (10 iterations) custom op: cuda - 32 threads per block / openmp - 12 threads

